

STA 3180 Statistical Modelling: Markov Chain Monte Carlo

Markov Chain Monte Carlo (MCMC)

Definition

Markov Chain Monte Carlo (MCMC) is a type of computational algorithm used to generate samples from a probability distribution. It is based on the concept of a Markov chain, which is a sequence of random variables that have a certain probability of transitioning from one state to another. MCMC algorithms are used in Bayesian inference and statistical modelling to approximate complex distributions and solve difficult problems.

Key Concepts

- **Markov Chain**: A Markov chain is a sequence of random variables that have a certain probability of transitioning from one state to another. The transition probabilities depend only on the current state and not on the previous states.
- **Markov Chain Monte Carlo (MCMC)**: MCMC is a type of computational algorithm used to generate samples from a probability distribution. It is based on the concept of a Markov chain.
- **Metropolis-Hastings Algorithm**: The Metropolis-Hastings algorithm is a popular MCMC algorithm. It is used to generate samples from a target distribution by constructing a Markov chain whose stationary distribution is the target distribution.
- **Gibbs Sampling**: Gibbs sampling is a special case of the Metropolis-Hastings algorithm. It is used to generate samples from a target distribution by constructing a Markov chain whose stationary distribution is the target distribution.

Coding Examples

Example 1: Metropolis-Hastings Algorithm

Start of Code

```
import numpy as np

def metropolis_hastings(target_dist, initial_state, num_samples):
    # Initialize the Markov chain
    chain = [initial_state]
    # Generate samples
    for i in range(num_samples):
        # Propose a new state
        proposed_state = propose_state(chain[-1])
```

```

    # Calculate the acceptance probability
    acceptance_prob = min(1, target_dist(proposed_state) /
                           target_dist(chain[-1]))
    # Sample a uniform random variable
    u = np.random.uniform()
    # Accept or reject the proposed state
    if u < acceptance_prob:
        chain.append(proposed_state)
    else:
        chain.append(chain[-1])

    return chain

```

End of Code

Example 2: Gibbs Sampling

Start of Code

```

import numpy as np

def gibbs_sampling(target_dist, initial_state, num_samples):
    # Initialize the Markov chain
    chain = [initial_state]
    # Generate samples
    for i in range(num_samples):
        # Sample a uniform random variable
        u = np.random.uniform()
        # Propose a new state
        proposed_state = propose_state(chain[-1], u)
        # Accept the proposed state
        chain.append(proposed_state)
    return chain

```

End of Code